



Big Data Processing (IE494)

Distributed Graph Processing using Apache Spark (Project Report)

Divyesh Ramani - 202201241

Manan Patel - 202201310

Prof. P.M. Jat

18 November, 2024

Problem and Solution

The project addresses the challenge of **Graph Processing at Scale**. As datasets grow in size and complexity, they can be represented as graphs consisting of nodes and edges. Querying such large-scale graphs on a centralized system becomes increasingly infeasible due to the limitations of memory and processing power.

To tackle this, we implemented a solution proposed by **Prof. Rajshekhar Sunderraman** and **Dr. Janani Balaji**, leveraging the in-memory processing capabilities and distributed architecture of **Apache Spark**. The approach involves decomposing a graph query into smaller components, enabling parallel processing across clusters. This method significantly enhances scalability and efficiency, making it well-suited for large-scale graph analysis.

Our implementation demonstrates how distributed systems like Apache Spark can effectively manage and process massive graphs, enabling advanced querying and insights on complex datasets.

Approach/Algorithm Implemented

1. Node and Edge File Parsing

- Objective: Read graph and query information from files and construct graph structures.
 - A. Node File Parsing:
 - Each line represents a node labeled as NodeLabel_NodeType.
 - The mapping of nodes to their types is stored in a dictionary.

```
FUNCTION read_nodes_from_file(file_path):
    INITIALIZE an empty dictionary `node`

    OPEN file at `file_path` in read mode with encoding 'utf-8-sig'
    FOR EACH line in the file:
        REMOVE leading and trailing whitespaces from `line`

        IF `line` is not empty AND `line` is not "Exit":
            IF "_" exists in `line`:
                SPLIT `line` into `node_label` and `node_type` at
                h"_"
                ADD `node_label` as key and `node_type` as value
                to `node`
            ELSE:
                PRINT "Skipping invalid line:" followed by `line`

    RETURN `node`
```

B. Edge File Parsing:

- Edges are represented as SourceNode_TargetNode_EdgeLabel.
- The edges are stored in dictionaries for both outbound ('o') and inbound ('i') connections in a temp_graph dictionary.

```
FUNCTION read_edges_from_file(file_path, temp_graph):
    OPEN file at `file_path` in read mode with encoding 'utf-8-sig'

    FOR EACH line in the file:
        REMOVE leading and trailing whitespaces from `line`

        IF `line` is not empty AND `line` is not "Exit":
            IF `line` starts with '\u0000':
                REMOVE the BOM character from `line`

            SPLIT `line` into `user_input` at "_"

            IF both `user_input[0]` and `user_input[1]` exist in
            `temp_graph`:
                # Handle outgoing edges
                IF edge label `(user_input[2], 'o')` exists in
                `temp_graph[user_input[0]][1]`:
                    APPEND `user_input[1]` (converted to integer) to
                    the list of values
                ELSE:
                    CREATE a new list with `user_input[1]` (converted
                    to integer) for the edge label

                # Handle incoming edges
                IF edge label `(user_input[2], 'i')` exists in
                `temp_graph[user_input[1]][1]`:
                    APPEND `user_input[0]` (converted to integer) to
                    the list of values
                ELSE:
                    CREATE a new list with `user_input[0]` (converted
                    to integer) for the edge label
```

2. Graph and Query Graph Construction

- Objective: Convert parsed nodes and edges into a distributed graph representation for both the main graph and the query graph.

A. Graph Representation:

- Nodes, their types, and edge details are organized into tuples and parallelized into RDDs (GraphRDD and QueryRDD).

B. Persistence:

- RDDs are persisted in memory for efficient reuse.

```

INITIALIZE an empty list `graph`

FOR EACH `key` in `temp_graph`:
    APPEND a tuple `(int(key), temp_graph[key][0], temp_graph[key][1])` to `graph`

CREATE `GraphRDD` by parallelizing `graph` using
`sc.parallelize(graph)`

PERSIST `GraphRDD` in memory for efficient reuse

INITIALIZE an empty list `query`

FOR EACH `key` in `temp_query`:
    APPEND a tuple `(int(key), temp_query[key][0], temp_query[key][1])` to `query`

CREATE `QueryRDD` by parallelizing `query` using
`sc.parallelize(query)`

PERSIST `QueryRDD` in memory for efficient reuse

```

3. Query Graph Segmentation

- Objective: Divide the query graph into segments for sequential processing.
 - A. Extract each node's label and associated edges.
 - B. Create segments:
 - For start and end nodes, include the node and all edges.
 - For intermediate nodes, include individual edges as separate segments.

```

SET `NumQuery` = the number of elements in `QueryRDD` (using
`QueryRDD.count()`)

INITIALIZE an empty list `segments`

SET `data` = the collected elements of `QueryRDD` (using
`QueryRDD.collect()`)

FOR `i` from 0 to `NumQuery` - 1:
    SET `node_label` = the label of the current query node
    (`data[i][1]`)
    SET `edge_list` = the keys of the edge dictionary for the current
    query node (`list(data[i][2].keys())`)

    INITIALIZE `tmp1` = a list containing `node_label`
    APPEND `edge_list` to `tmp1`

    INITIALIZE `tmp2` = a list containing `node_label`
    FOR EACH `edge` in `edge_list`:
        APPEND `[edge]` to `tmp2` (single-element list)

    IF `i` is the first node (`i == 0`) OR the last node
    (`i == NumQuery - 1`):
        APPEND `tmp1` to `segments`
    ELSE:
        APPEND `tmp2` to `segments`

```

4. Valid Candidate Initialization

- Objective: Initialize the search space with the entire graph.
 - A. Create validRDD, a transformation of GraphRDD that pairs each node with an initially empty list for path tracking.

5. Iterative Subgraph Matching

- Objective: Iteratively refine the search space to match the query graph structure.
 1. Node Label Filtering:
 - Filter nodes in validRDD that match the current query segment's node label.
 - B. Edge Filtering:
 - Further filter based on edges (label and direction) in the query segment.
 - C. Broadcast Shortlisted Nodes:
 - Broadcast the filtered results to improve parallelism in the next steps.
 - D. Neighbor Expansion:
 - Explore neighbors of shortlisted nodes using their edges.
 - Add paths of matched nodes to track query subgraph matches.

```
FOR `i` from 0 to `NumQuery - 1`:  
  # Step 1: Extract the node label and edges from the current query  
  segment  
  SET `node_label` = the first element of `segments[i]` (current  
  query node label)  
  SET `edge_list` = all elements of `segments[i]` starting from the  
  second (current query edges)
```

```

# Step 2: Filter valid nodes by node label
SET `shortlistedRDD` = `validRDD` filtered to only include nodes
with label `node_label`

# Step 3: Extract edge labels and directions from the edge list
INITIALIZE empty lists `e_lab` and `e_dir`
FOR EACH `edge` in `edge_list`:
    APPEND the edge label (`edge[0][0]`) to `e_lab`
    APPEND the edge direction (`edge[0][1]`) to `e_dir`

# Step 4: Further filter shortlisted nodes based on edges
FOR `j` from 0 to `len(e_lab) - 1`:
    UPDATE `shortlistedRDD` to only include nodes with
    `(e_lab[j], e_dir[j])` as keys in their edge dictionary

# Step 5: Collect shortlisted nodes and broadcast them
SET `shortlisted_data` = collected elements of `shortlistedRDD`
BROADCAST `shortlisted_data` as `shortlisted_broadcast`

# Step 6: Create a new RDD for next possible nodes (neighbors)
INITIALIZE `newRDD` as an empty RDD
FOR EACH `node` in `shortlisted_data`:
    SET `current_label` = the ID of the current node
    (`node[0][0]`)
    FOR `k` from 0 to `len(e_lab) - 1`:
        UNION the result of filtering `GraphRDD` for neighbors
        reachable via `(e_lab[k], e_dir[k])`:
            - Filter nodes where `x[0]` exists in the adjacency
              list `(e_lab[k], e_dir[k])` of `node`
            - Map the filtered nodes to `[x,
              path_with_current_label]`, where
              `path_with_current_label` is the previous path
              extended with `current_label`

# Step 7: Persist the new RDD
PERSIST `newRDD` to optimize further computations

# Step 8: Update `validRDD` for the next iteration
SET `validRDD` = `newRDD`

```

6. Result Collection

- Objective: Collect and print all matched subgraphs.
 - A. Final validRDD contains potential subgraph matches.
 - B. Extract and print paths from matched nodes.

Python Code

[Google Colab: Distributed Graph Processing Using Apache Spark](#)

Testing

1. Test Cases for `read_nodes_from_file`

- **Objective:** Verify that the function correctly parses nodes from the file and handles malformed data.
- **Test Inputs:**
 - A. A file with valid node entries (NodeID_NodeLabel).
 - B. A file with invalid entries.
 - C. A file with a mix of valid and invalid lines.
- **Expected Outputs:**
 - A. The function returns a dictionary with valid node-label mappings.
 - B. It skips invalid lines and prints appropriate warnings.
- **Test Scenarios:**
 - A. Empty file should return an empty dictionary.
 - B. File with only invalid lines should return an empty dictionary with warnings.

2. Test Cases for `read_edges_from_file`

- **Objective:** Verify that edges are added correctly to `temp_graph` and handle invalid entries gracefully.
- **Test Inputs:**
 - A. A file with valid edge entries (FromNodeID_ToNodeID_EdgeLabel).
 - B. A file with missing nodes (e.g., 3_4_EdgeLabel where 3 or 4 is not in `temp_graph`).
 - C. A file with invalid formats (e.g., 1-2-EdgeLabel).
- **Expected Outputs:**
 - A. Valid edges are added correctly to the graph.
 - B. Missing nodes are skipped, and warnings are generated.
- **Test Scenarios:**
 - A. In case of valid file, graph edges should update as expected.
 - B. In case of empty file graph remains unchanged.
 - C. In case of file with invalid edges graph remains unchanged with appropriate warnings.

3. Test Cases for Graph Construction

- **Objective:** Validate that graph and query lists are constructed correctly from `temp_graph` and `temp_query`.
- **Test Inputs:**
 - A. Manually created `temp_graph` and `temp_query` dictionaries.


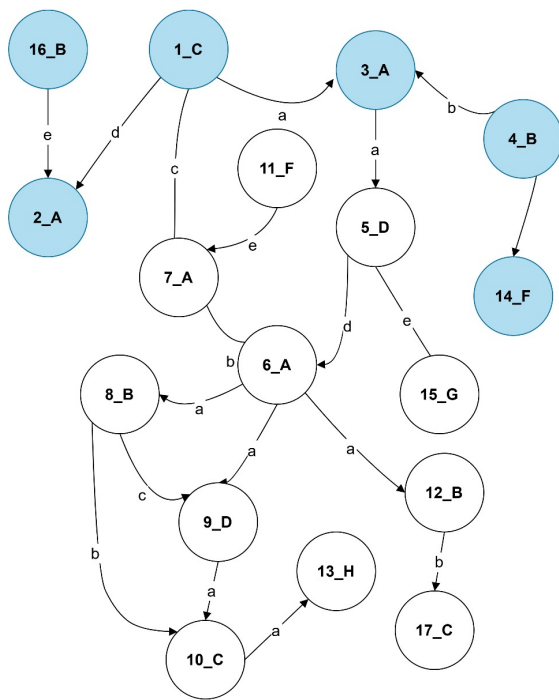
- **Expected Outputs:**
 - A. Correct tuples are created for all nodes in temp_graph and temp_query.
- **Test Scenarios:**
 - A. temp_graph with no nodes should return an empty RDD.
 - B. temp_graph with multiple nodes should return an RDD with the correct structure.

4. Test Cases for Query Segmentation

- **Objective:** Verify the correctness of the segmentation logic.
- **Test Inputs:**
 - A. A QueryRDD with varying sizes (small, medium, large).
 - B. Query nodes with multiple and zero edges.
- **Expected Outputs:**
 - A. segments list correctly segments nodes and edges.
- **Test Scenarios:**
 - A. Single query node with no edges should result in one segment with no edge information.
 - B. Multiple query nodes with edges should result in segments correctly split nodes and edges.

5. Test Cases for Matching Logic

- **Objective:** Validate the iterative node matching and edge filtering.
- **Test Inputs:**
 - A. Simple graphs with known matches.
 - B. Complex graphs with cyclic dependencies.
- **Expected Outputs:**
 - A. Correctly shortlisted nodes after each iteration.
 - B. Accurate final validRDD with matched nodes and paths.
- **Test Scenarios:**
 - A. Simple graph with direct matches → Final RDD contains expected matches.
 - B. Graph with no matches → Final RDD is empty.
 - C. Graph with multiple valid paths → All paths are identified.



```

1_B
2_A
3_C
4_A
5_B
6_F
Exit

```

Query Node

```

1_2_e
3_2_d
3_4_a
5_4_b
5_6_c
Exit

```

Query Edge

```

7_A
8_B
9_D
10_C
11_F
12_B
13_H
14_F
15_G
16_B
17_C
Exit

```

Graph Nodes

```

7_6_d
6_8_a
6_9_a
8_9_c
8_10_b
9_10_a
11_7_e
6_12_a
10_13_a
4_14_c
5_15_e
16_2_e
12_17_b
Exit

```

Graph Edges

```
[16, 2, 1, 3, 4, 14]
```

Output using the python code

References

Balaji, Janani, and Rajshekhar Sunderraman. "Distributed graph path queries using spark." *2016 IEEE 40th Annual Computer Software and Applications Conference (COMPSAC)*. Vol. 2. IEEE, 2016.